


# Query Processing and Optimization in Distributed Databases




Erdem USLU

Nihal SARMAŞIK

# CONTENT

- Introduction to Distributed DBMS
- Storing Data in Distributed DBMS
- Distributed Query Processing
- Cost-Based Query Optimization



# Introduction to Distributed DBMS

## Main Considerations:

### ➤ Distributed Transaction Atomicity

- ❖ users should be able to write transactions that access and update data at several sites just as they would write transactions over purely local data

### ➤ Location Transparency


- ❖ user does not have to know the location of the data
- ❖ data requests automatically forwarded to appropriate sites

# Introduction to Distributed DBMS

## Types of Distributed Database System

### ➤ Homogeneous Distributed Database System

- ❖ All sites of the database system have identical setup, i.e., same database system software.
- ❖ The underlying operating system may be different




# Introduction to Distributed DBMS

## Types of Distributed Database System

### ➤ Heterogeneous Distributed Database System (Multidatabase System)

Each site may run different database system/software



# Storing Data in Distributed DBMS

## ➤ Fragmentation


- ❖ Division of relation  $r$  into fragments  $r_1, r_2, \dots, r_n$  which contain sufficient information to reconstruct relation  $r$
- ❖ allows a relation to be split so that tuples are located where they are most frequently accessed

## ➤ Fragmentation types

- ❖ Horizontal Fragmentation
- ❖ Vertical Fragmentation
- ❖ Mixed Fragmentation

# Storing Data in Distributed DBMS

## Horizontal Fragmentation ( Row-wise )



$$R = R1 \cup R2$$


# Storing Data in Distributed DBMS

## Vertical Fragmentation (Column-wise )

❖ The schema for relation  $r$  is split into several smaller schemas

❖ All schemas must contain a common candidate key (or superkey) to ensure lossless join property

❖ A special attribute, the tuple-id attribute may be added to each schema to serve as a candidate key



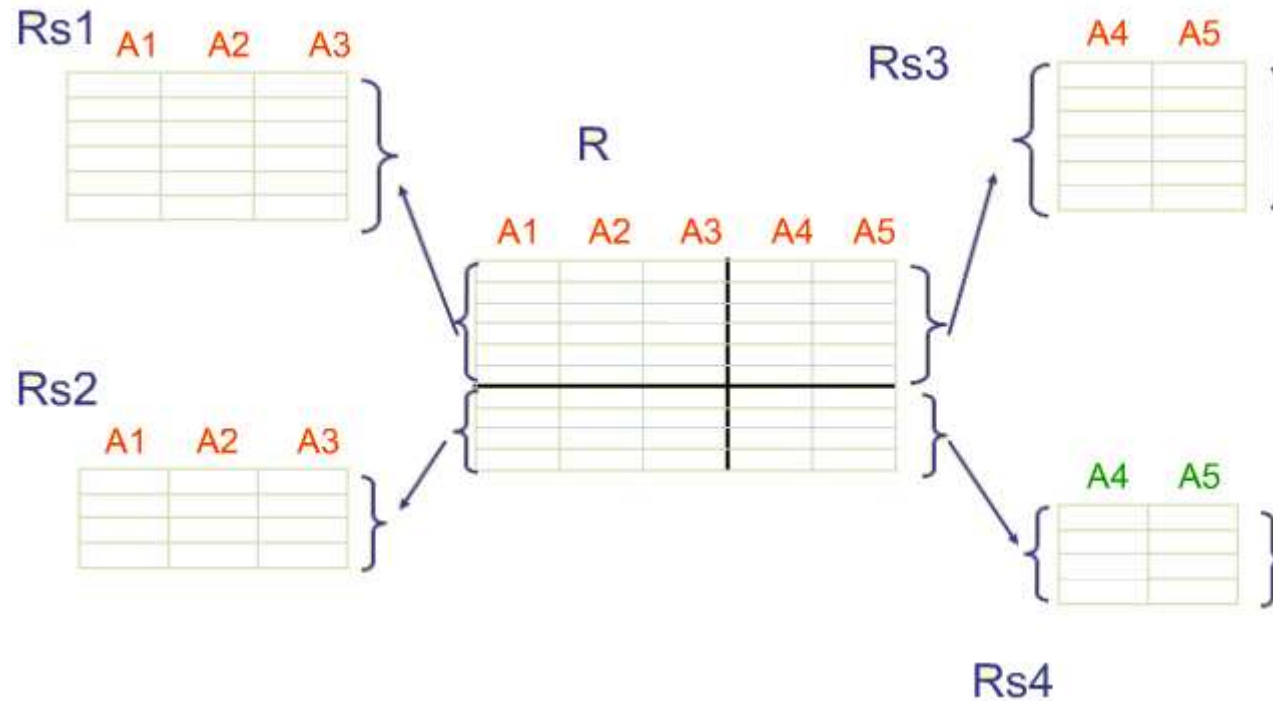
$$R = R1 \bowtie R2$$

$$R1.TID = R2.TID$$



# Storing Data in Distributed DBMS

## Mixed fragmentation




$$R = (Rs1 \cup Rs2) \bowtie (Rs3 \cup Rs4)$$

# Storing Data in Distributed DBMS

## ➤ Replication

- ❖ Multiple copies of data, stored in different sites



## ➤ Advantages of Replication

- ❖ **Availability:** failure of site containing relation  $r$  does not result in unavailability of  $r$  if replicas exist
- ❖ **Parallelism:** queries on  $r$  may be processed by several nodes in parallel.
- ❖ **Reduced data transfer:** relation  $r$  is available locally at each site containing a replica of  $r$

# Distributed Query Processing


- Criteria for measuring the cost of a query evaluation strategy
  - ❖ For centralized DBMSs number of disk accesses (# blocks read / written)
  - ❖ For distributed databases, additionally
    - ✓ The cost of data transmission over the network
    - ✓ Potential gain in performance from having several sites processing parts of the query in parallel

# Distributed Query Processing


## Nonjoin Queries in Distributed DBMS

Example:

$EMP = EMP1 \cup EMP2$  (Horizontal Fragmentation)



$EMP1 = \sigma_{emp\_id < 5000}(EMP)$



$EMP2 = \sigma_{emp\_id \geq 5000}(EMP)$

# Distributed Query Processing

## Nonjoin Queries in Distributed DBMS

**QUERY:** Select all employees whose ages between 20 and 30.



```
SELECT emp_id, name  
FROM EMP1
```

```
WHERE age > 20 AND age < 30
```

UNION

```
SELECT emp_id, name  
FROM EMP2
```

```
WHERE age > 20 AND age < 30
```

# Distributed Query Processing

## Nonjoin Queries in Distributed DBMS

**QUERY:** What is average age of all employee?



```
SELECT      (T1.age+T2.age) /  
(T1.emp_number+T2.emp_number)
```

```
FROM
```

```
(SELECT SUM(age) as age, count(*) as  
emp_number FROM EMP1) as T1 ,
```

```
(SELECT SUM(age) as age, count(*) as  
emp_number FROM EMP2) as T2
```

# Distributed Query Processing

## Join Queries in Distributed DBMS

### ➤ JOIN STRATEGY

- ❖ Ship whole
- ❖ Fetch as needed
- ❖ Semijoins
- ❖ Bloomjoins

Which strategy is better for me?



# Distributed Query Processing

## Join Queries in Distributed DBMS (Ship Whole)

**Ship Whole:** Transferring the complete relation

**Example:**

R	A	B
	3	7
	1	1
	4	6
	7	7
	4	5
	6	2
	5	7

S	B	C	D
	9	8	8
	1	5	1
	9	4	2
	4	3	3
	4	2	6
	5	7	8

$R \bowtie S$	A	B	C	D
	1	1	5	1
	4	5	7	8

**QUERY:** The query asks for  $R \bowtie S$



# Distributed Query Processing

## Join Queries in Distributed DBMS (Ship whole)

### COST ANALYSIS

#### ➤ When execution at nodeR

- ❖ nodeR: send data request message (relation S) to nodeS
- ❖ nodeS: send requested data (relation S) to nodeR

**Total costs: 2 messages, 18 attribute value**

#### ➤ When execution at nodeS

- ❖ nodeS: send data request message (relation R) to nodeR
- ❖ nodeR: send requested data (relation R) to nodeS

**Total costs: 2 messages, 14 attribute value**

# Distributed Query Processing

## Join Queries in Distributed DBMS (Fetch As Needed)

**Fetch As Needed:** Transferring the relation piecewise

**Example:**

R

A	B
3	7
1	1
4	6
7	7
4	5
6	2
5	7

S

B	C	D
9	8	8
1	5	1
9	4	2
4	3	3
4	2	6
5	7	8

$R \bowtie S$

A	B	C	D
1	1	5	1
4	5	7	8

**QUERY:** The query asks for  $R \bowtie S$

# Distributed Query Processing

## Join Queries in Distributed DBMS (Fetch As Needed)

### COST ANALYSIS

#### ➤ When execution at nodeR

- ❖ **nodeR**: send data request message (tuples of relation S with B = '7') to nodeS
- ❖ **nodeS**: send requested data (0 tuples of relation S with B = '7') to nodeR
- ❖ **nodeR**: send data request message (tuples of relation S with B = '1') to nodeS
- ❖ **nodeS**: send requested data (1 tuple of relation S with B = '1') to nodeR
- ...

**Total costs:  $7 * 2 = 14$  messages,  $7 + 2 * 3 = 13$  attribute value**

#### ➤ When execution at nodeS

- ❖ **nodeS**: send data request message (tuples of relation S with B = '9') to nodeR
- ❖ **nodeR**: send requested data (0 tuples of relation S with B = '9') to nodeS
- ❖ **nodeS**: send data request message (tuples of relation S with B = '1') to nodeR
- ❖ **nodeR**: send requested data (1 tuple of relation S with B = '1') to nodeS
- ...

**Total costs:  $6 * 2 = 12$  messages,  $6 + 2 * 2 = 10$  attribute value**

# Distributed Query Processing

## Join Queries in Distributed DBMS

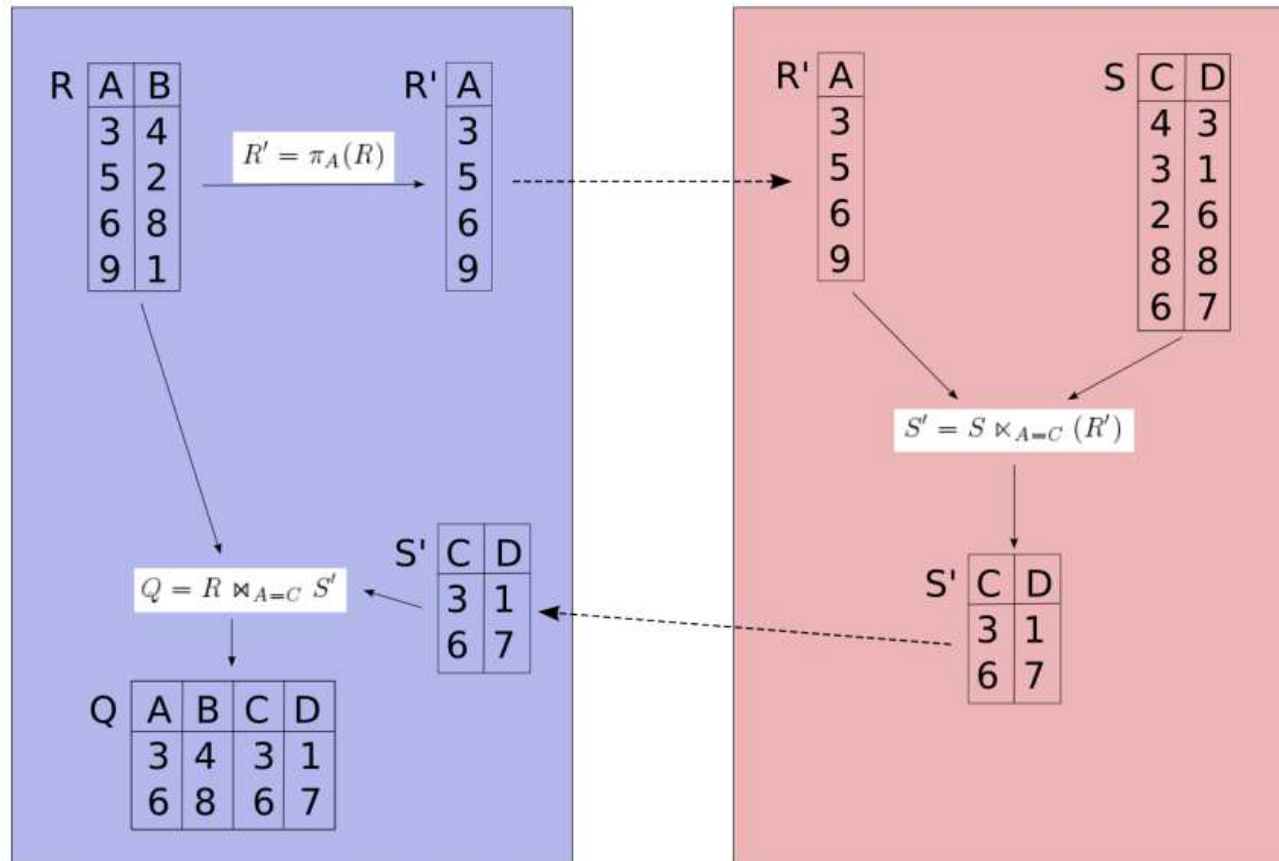
### **Ship Whole vs Fetch as needed:**

- Fetch as needed results in  
a high number of messages
- Ship whole results in  
high amounts of transferred data

# Distributed Query Processing

## Join Queries in Distributed DBMS (Semijoin)

**Semijoin:** Requesting all join partners in just one step



# Distributed Query Processing

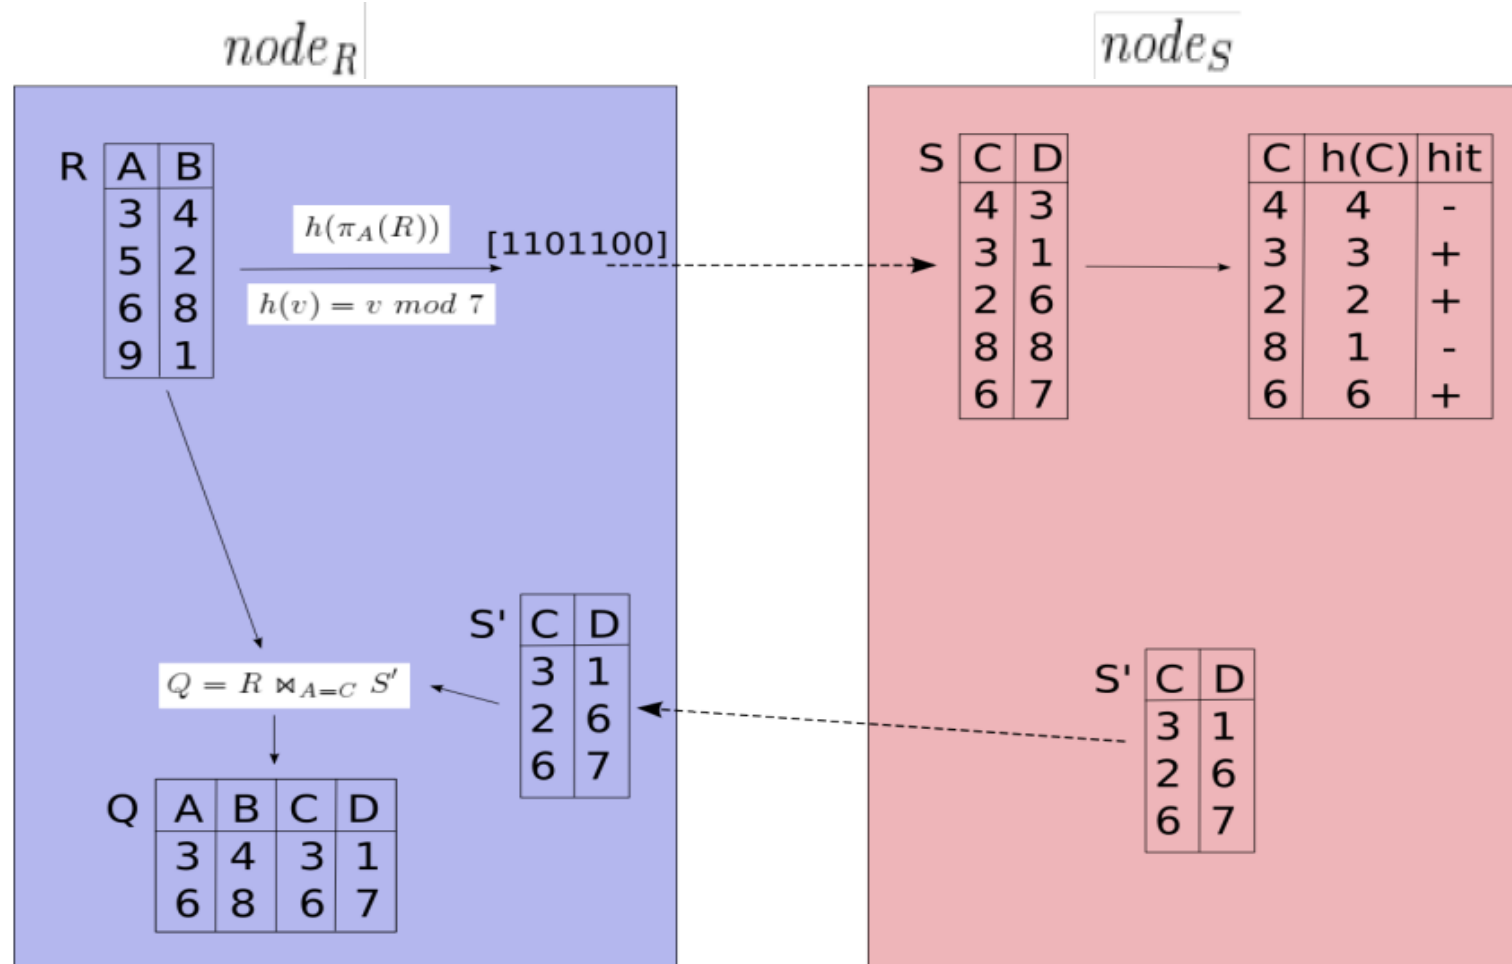
## Join Queries in Distributed DBMS (Bloom join)

### Bloom join:

- Also known as bit-vector join
- Avoiding to transfer all join attribute values to the other node
- Instead transferring a bitvector  $B[1 : : n]$
- Transformation
  - ❖ Choose an appropriate hash function  $h$
  - ❖ Apply  $h$  to transform attribute values to the range  $[1 : : n]$
  - ❖ Set the corresponding bits in the bitvector  $B[1 : : n]$  to 1

# Distributed Query Processing

## Join Queries in Distributed DBMS (Bloom join)



# Distributed Query Processing

## Join Queries in Distributed DBMS (Bloom join)

### Conclusion:

- Transferring the bit-vector reduces network load
- Bit-vector only indicates potential join partners because multiple attribute values might map to the same hash value  
*Might result in transferring unnecessary tuples*
- Requirements: an appropriate hash function  $h$  and  $n$  needs to be large enough to avoid a high number of collision



# Distributed Query Processing

## Join Queries in Distributed DBMS

### **Semijoin vs bloom-join**

- \* The cost of shipping is less in bloom-join since bit-vector is used rather than projection
- \* In Bloom join the size of the reductioned part which transferred back is likely to be larger than in Semijoin, so the costs of shipping are likely to be higher

# Cost-Based Query Optimization

## Main Consideration for Query Optimization

- Communication cost
- If there is several copies of a relation, decide which copy to use
- Amount of data being shipped
- Relative processing speed at each site
- Site selection

# Cost-Based Query Optimization

## ➤ Global plan:

- ❖ Includes several local plans (subqueries)
- ❖ If response time is critical, subqueries can be carried out in parallel
- ❖ Local plans constructed by optimizer of each site

# THANK YOU!

